

AC290: Extreme Computing

Intro to GPU Computing

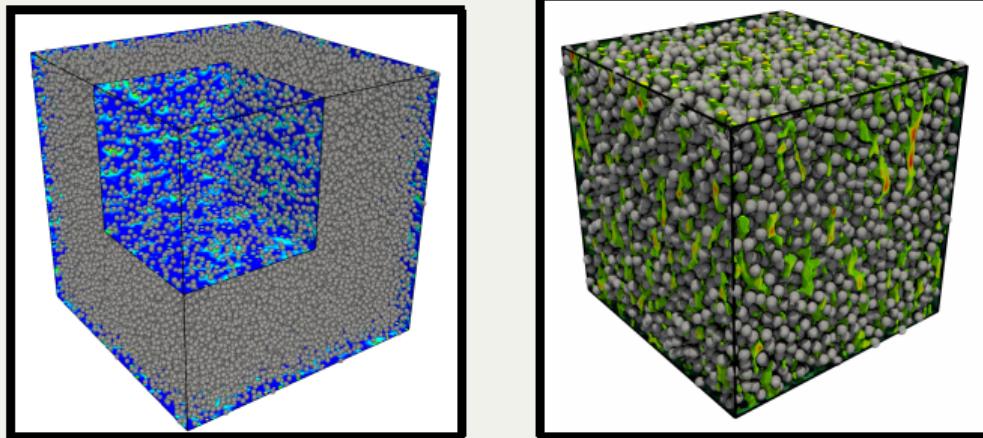
Daniel Willen

The MITRE Corporation

Lecture 1 Recap

- ▶ Compiling and running GPU code on a cluster
- ▶ Cuda kernel configuration and thread management
- ▶ Memory management

Any questions? (save HW-specific questions for the end)



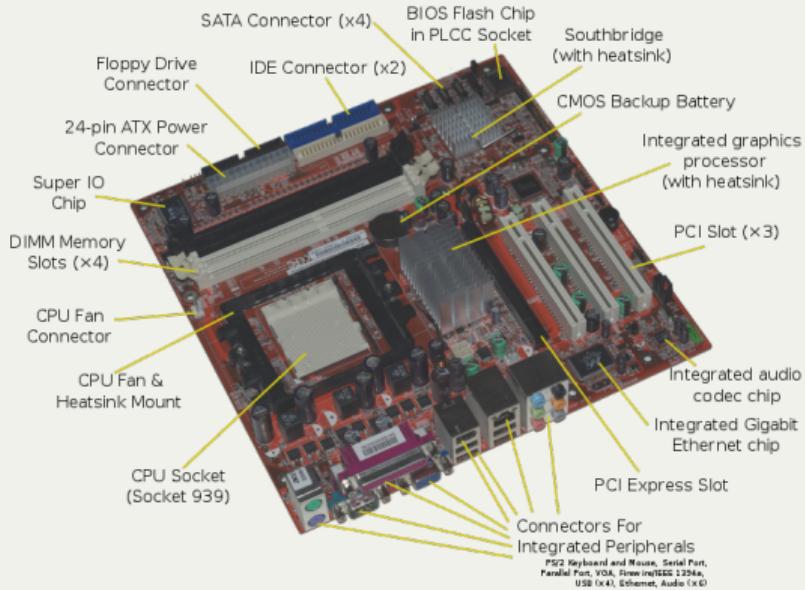
Overview

1. GPU Architecture
2. GPU Memory Types
3. Cuda-aware MPI
4. Homework Help

Objectives:

1. Become familiar with GPU architecture terms
2. Become familiar with what intermediate and advanced Cuda constructs exist

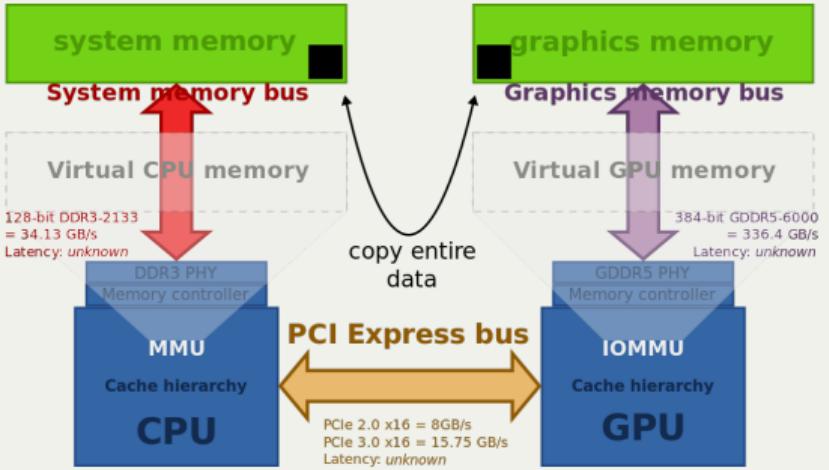
GPUs and Your Computer



- ▶ Motherboard connects most of the components in a computer
 - ▶ CPU, RAM, GPU, sound card, ...
- ▶ CPU and GPU are connected via the Peripheral Component Interface Express bus
 - ▶ PCI Express or PCIe
 - ▶ Moving data from one to the other requires communication!

PCIe Bus

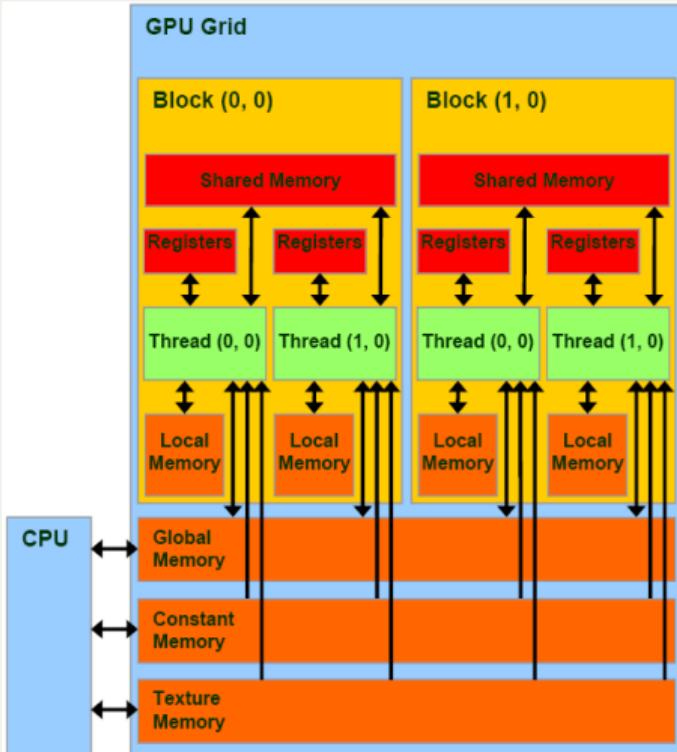
- ▶ Processors are fast:
 - ▶ NVIDIA V100: 900 GB/s
 - ▶ Intel Cascade Lake: 128 GB/s
- ▶ CPU-GPU communication is slow:
 - ▶ PCIe Gen3: 32 GB/s
 - ▶ Infiniband HDR: 50 GB/s
 - ▶ NVLink: 300 GB/s
- ▶ Communication is a bottleneck!
 - ▶ Avoid it (algorithmic, GPU-centric)
 - ▶ Hide it (async)



Memory Layout

Several different types of GPU memory...

- ▶ Red – on-chip
- ▶ Orange – off-chip
- ▶ cudaMemcpy copies to global memory
- ▶ Other types need to be set explicitly



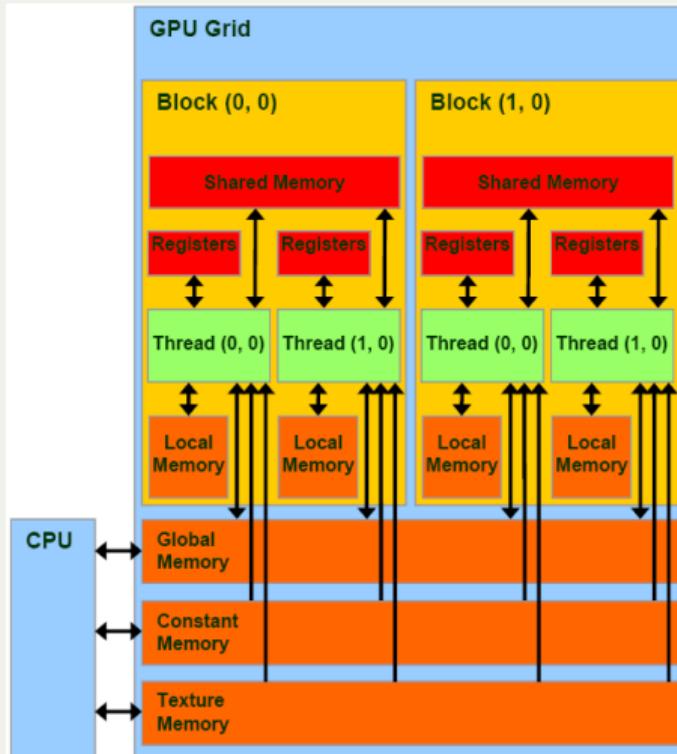
Nvidia V100 Architecture

- ▶ Streaming Multiprocessor (SM): 64 FP32 cores, 64 INT32 cores, 32 FP64 cores



- ▶ Recall:
 - ▶ Threads are organized into thread blocks
 - ▶ Thread blocks are organized into a grid
- ▶ Compiler allocates thread blocks to SMs (32 FP64 cores/SM)
- ▶ Thread blocks are allocated to ‘warps’ of threads (32/warp)
- ▶ There are hardware limits on...
 1. Threads per warp (32) (threads can be executed in half-warps)
 2. Threads per block
 3. Active blocks
 4. Dimension of thread blocks
 5. ... and many more
- ▶ This is important for optimizations
- ▶ Further reading: see References

Memory Layout



1. Register: On-chip, per-thread
2. Shared: On-chip, all threads in block
3. Constant: Off-chip, always available, read-only
4. Texture: Read-only, interpolation!
5. Local: Off-chip, per-thread
6. Global: Off-chip, always available, default

Memory Types

10/26

- ▶ Access latency (in clock cycles) can depend on the type of access
 - ▶ If all threads access the same memory address, the value can be broadcast
 - ▶ If not, accesses are serialized
- ▶ Penalty – roughly how much longer an access takes compared to register

Type	Location	Lifetime	Penalty	Latency (cycles)	Size
Register (L1)	On-chip	Thread	1×	28	up to 128KB/SM
Shared	On-chip	Block	1×	20-80	up to 96KB/SM
Constant	Off-chip	App	1×	10-300	2KB
Local (L2)	Off-chip	Thread	100×	193	6124KB
Global	Off-chip	App	100×	~ 235	~ 16-32GB

Device Constant Memory

11/26

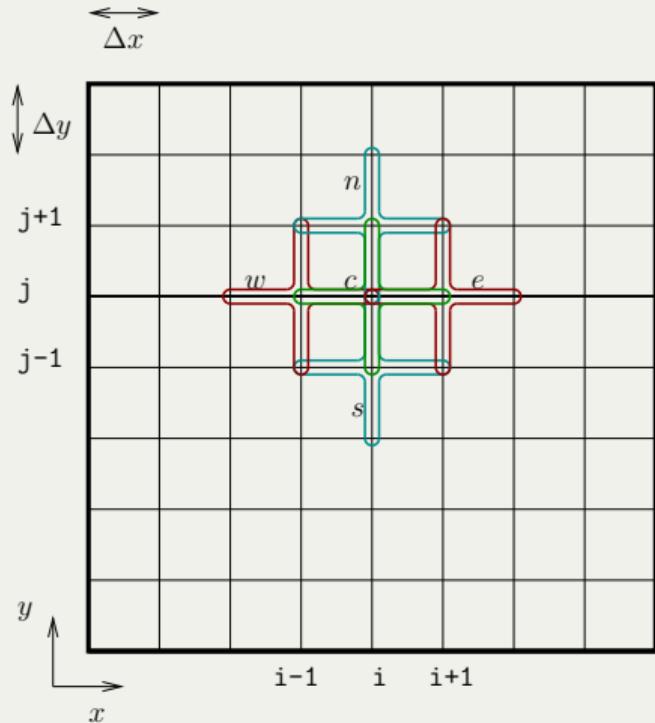
```
/* Set up device constant memory */
int var = 1;           // Host
__constant__ int _var; // Device Constant
cudaMemcpyToSymbol(_var, &var, sizeof(int));

/* Use in a kernel */
__global__ void k1(void) {
    int twice_var = 2*_var;
}
```

- ▶ `cudaMemcpyToSymbol(`
 `const void* symbol,`
 `const void* src, ...)`
- ▶ ‘symbol is a device symbol
address and is a variable
that resides in global or
constant memory space’

Shared Memory

12/26

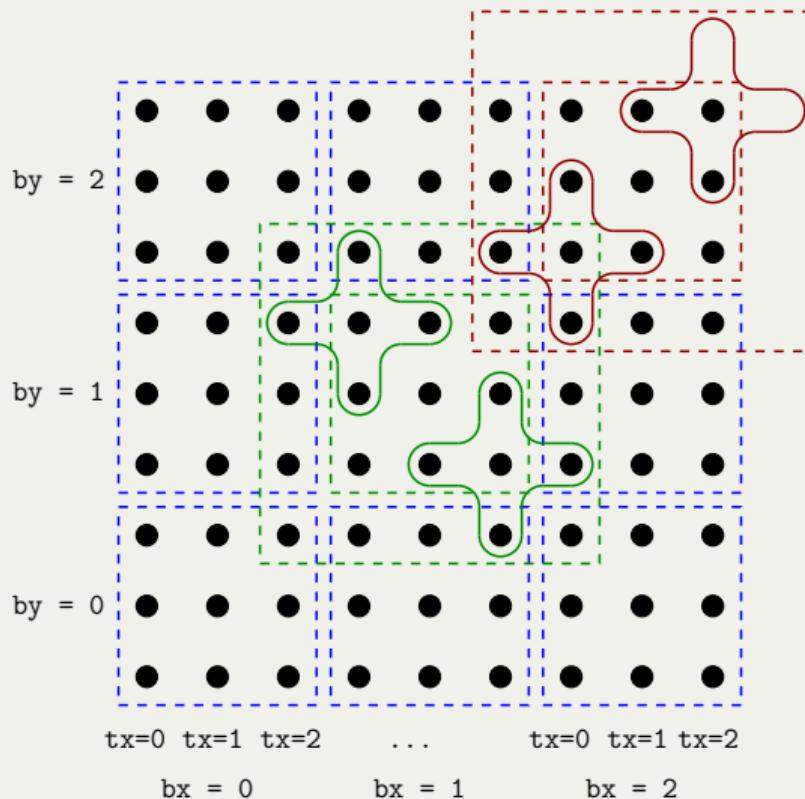


Finite-difference stencils reuse grid points:

- ▶ Central difference: $5 \times$ in 2-D, $7 \times$ in 3-D
- ▶ Each thread reads the same value from global memory to register. This is slow!
- ▶ Solution: Store the value in shared memory that can be accessed by all threads in a block
 - ▶ Read once, use many times
 - ▶ Should not write back to shared memory
 - write elsewhere!

Shared Memory

13/26



- ▶ Each thread block needs data from one more thread on each size
- ▶ Thus, each thread block size is 2 greater in each dimension than the number of points to compute
- ▶ We need to take care not to calculate the value at the edge of the stencil!

Shared Memory Example

14/26

```
// Set up kernel execution. Assume we're using the max number of threads in
// each dimension
int threads_x = MAX_THREADS_DIM;
int threads_y = MAX_THREADS_DIM;

// Each block has 2 extra threads in our example
int blocks_x = ceil(nx / (threads_x - 2)); // Should cast division to double!
int blocks_y = ceil(ny / (threads_y - 2)); // Should cast division to double!

dim3 num_blocks(blocks_x, blocks_y);
dim3 dim_blocks(threads_x, threads_y);

// Launch the kernel as usual
my_kernel<<<num_blocks, dim_blocks>>>(_u, nx, ny);
```

Shared Memory Example

15/26

```
my_kernel(double *u, int nx, int ny) {
    // Declare shared memory
    __shared__ double s_u[MAX_THREADS_DIM*MAX_THREADS_DIM];

    // Shared memory index in shared memory grid
    int si = threadIdx.x;    int sj = threadIdx.y;
    int sc = si + sj*blockDim.x;

    // Global memory index in regular grid
    // Subtraction accounts for shared memory overlap
    int ti = blockIdx.x*blockDim.x + threadIdx.x - 2*blockIdx.x;
    int tj = blockIdx.y*blockDim.y + threadIdx.y - 2*blockIdx.y;
    int cc = ti + nx*tj;

    // Load shared memory
    if (ti < nx && tj < ny) { // Only load data in computational grid
        s_u[sc] = u[cc];           // Shared memory <- global memory
    }
}
```

Shared Memory Example

16/26

```
// Ensure initial load is done
__syncthreads();

// Perform Jacobi; avoid shared memory 'shell'
if ((ti >= 1 && ti < (nx - 1)) &&          // Inside global domain (x)
    (tj >= 1 && tj < (ny - 1)) &&          // Inside global domain (y)
    (si > 0 && si < (blockDim.x - 1)) && // Off shared memory shell (x)
    (sj > 0 && sj < (blockDim.y - 1))) { // Off shared memory shell (y)

    // Center: u[i + j*nx] = s_u[si + sj*blockDim.x]
    // East:   u[i+1 + j*nx] = s_u[si+1 + sj*blockDim.x]
    // North:  u[i + (j+1)*nx] = s_u[si + (sj+1)*blockDim.x]
}

} // end of kernel
```

Harder, Better, Faster, Stronger

17/26

- ▶ Most HPC applications want to run larger simulations
 - ▶ Weak scaling: domain size scales with number of processors
 - ▶ Computational work per processor is constant
- ▶ Or, run faster
 - ▶ Strong scaling: Domain size is constant
 - ▶ Computational work per processor decreases
- ▶ High-end GPUs only have 32GB of device memory
 - ▶ Fluids have four fields: u, v, w, p
 - ▶ `sizeof(double)` = 8 bytes
 - ▶ Total domain memory = 32GB = (4 fields) $\left(\frac{(8n_x)^3 \text{bytes}}{\text{field}} \right)$
 - ▶ Domain size: $\rightarrow n_x \approx 250$: This is small!
 - ▶ Current state-of-the-art is $10,000^3$

Communication Breakdown

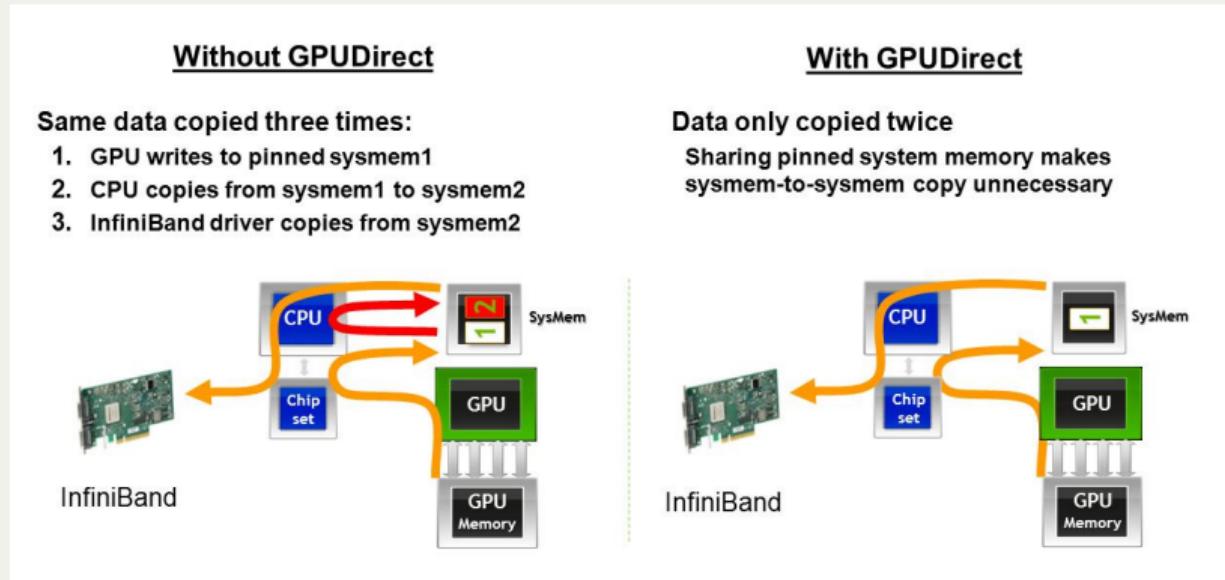
18/26

- ▶ In order to run larger simulations, we need to include more GPUs
- ▶ This results in memory distributed among the GPUs
- ▶ This requires communication for:
 - ▶ Boundary condition updates at new time step
 - ▶ Moving objects from one subdomain to another (e.g., particles in a fluid)
 - ▶ Global reductions (scalar products, min and max, sums, ...)
 - ▶ Most systems of linear equations are solved iteratively
 - ▶ Several reductions per iteration!
- ▶ Recall that communication to/from the GPU is $O(10)$ slower than the bandwidth!

- ▶ MPI defines standard by which programs running in distributed memory spaces can communicate
- ▶ Originally written for CPU-CPU communication.
- ▶ Old GPU-GPU communication:
 1. Copy memory from GPU to CPU
 2. MPI: Copy memory from CPU to CPU (inter- or intra-compute node)
 3. Copy memory from CPU to GPU
- ▶ Recent (2013) advances allow MPI to recognize GPU memory spaces
 - ▶ Abstracts and optimizes memory copying
 - ▶ Allows for transparent use of communication acceleration technologies
- ▶ Nvidia is very aware that inter-GPU communication is an issue
 - ▶ GPUDirect
 - ▶ NVLink and NVSwitch
 - ▶ Many-GPU implementation is *not* necessarily inefficient!

GPUDirect (2010)

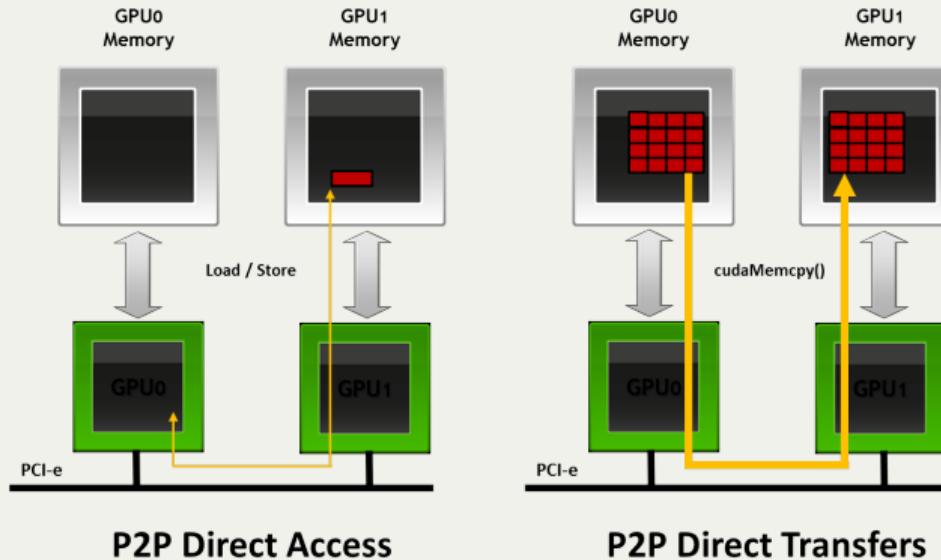
- ▶ Inter-node
- ▶ `cudaMallocHost` uses this



GPUDirect

GPUDirect Peer-to-Peer (P2P) (2011)

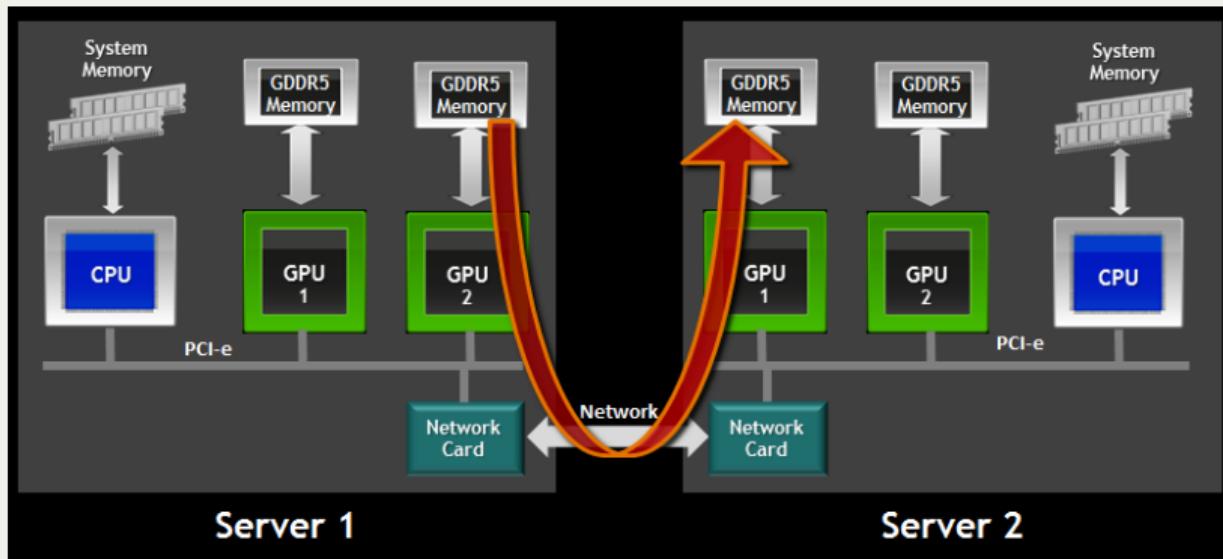
- ▶ Intra-node
- ▶ `cudaMemcpy(..., ..., ..., cudaMemcpyDeviceToDevice)`



GPUDirect

GPUDirect Remote-Direct Memory Access (RDMA) (2012)

- ▶ Inter-node
- ▶ Exposed via MPI



Cuda-Aware MPI Example

23/26

```
// Init MPI
MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &nprocs);

// Init data
int local = rank; // Send this
int remote = -1; // Recv here

int *_local, *_remote;
cudaMalloc(&_local, sizeof(int));
cudaMalloc(&_remote, sizeof(int));
// Copy host->device
cudaMemcpy(_local, local, sizeof(int), ...);
cudaMemcpy(_remote, remote, sizeof(int), ...);
```

- ▶ Set up MPI as usual
- ▶ Local data set to rank number
- ▶ Remote data set to -1
- ▶ Allocate and copy memory to GPU

Cuda-Aware MPI Example

24/26

```
MPI_win win;  
MPI_Win_create(_remote, 1*sizeof(int),  
               sizeof(int), MPI_INFO_NULL,  
               MPI_COMM_WORLD, &win);  
  
MPI_Win_fence(0, win);  
  
int target = (rank + 1) % nprocs;  
MPI_Put(_local, 1, MPI_INT,  
        target, 0, 1, MPI_INT, win);  
  
MPI_Win_fence(0, win);  
  
cudaMemcpy(local, _local, sizeof(int));  
cudaMemcpy(remote, _remote, sizeof(int));
```

- ▶ MPI_win is one-sided
 - ▶ Does not require MPI_Send or MPI_Recv
 - ▶ Processes put data directly in remote processes' memory (RDMA)
- ▶ Need to declare which memory is accessible
- ▶ Need to declare when it is accessible
- ▶ Put my info in the target's window
- ▶ Use GPU memory exactly like I would CPU
 - ▶ Note that `_remote` and `_local` are passed to the MPI calls!

LBM & GPUs

The basic LBM scheme reads

$$f_p(x + c_p, t + 1) = f_p(x, t) + \omega[f_p^{\text{eq}}(\rho, \mathbf{u}) - f_p](x, t),$$

and does not look too different from a matrix problem.

In fact, let's consider a rectangular mesh. Then,

$$f_p^{(n+1)}[i', j'] = f_p^{(n)}[i, j] + \omega [f_p^{\text{eq}}[i, j] - f_p[i, j]]$$

where i' and j' are destination (post-streaming) nodes.

Beyond similarity, what really matters is to take into account *data locality* that is critical for efficient GPU computing.

References

1. <https://devblogs.nvidia.com/inside-volta/>
2. <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#programming-model>
3. <https://docs.nvidia.com/cuda/parallel-thread-execution/index.html>
4. <http://www.ce.jhu.edu/dalrymple/classes/602/Class13.pdf>
5. <https://cs.nyu.edu/courses/spring12/CSCI-GA.3033-012/lecture5.pdf>
6. <https://www.anandtech.com/show/3809/nvidias-geforce-gtx-460-the-200-king/2>
7. <https://devblogs.nvidia.com>
8. https://www.nvidia.com/content/cudazone/download/Getting_Started_w_CUDA_Training_NVISION08.pdf
9. <https://stackoverflow.com/questions/3606636/cuda-model-what-is-warp-size>
10. <https://stackoverflow.com/questions/10460742/how-do-cuda-blocks-warps-threads-map-onto-cuda-cores>
11. <https://docs.nvidia.com/cuda/cuda-c-best-practices-guide/index.html#occupancy>
12. Dissecting the NVIDIA Volta GPU Architecture via Microbenchmarking